

**UNIVERSITY OF WATERLOO**  
**Faculty of Engineering**

**Backing Database System for a Messaging Platform Startup**

Ten Thousand Coffees  
Toronto, ON

Clement Hoang  
2B Software Engineering  
20531116  
c8hoang  
May 4th, 2015

Clement Hoang  
333 King Street North  
Waterloo, ON. 2Z1 N2J

May 11, 2015

Andrew Morton, Director  
Software Engineering Program  
University of Waterloo  
200 University Ave. West  
Waterloo, Ontario. N2L 3G1

Dear Doctor Morton:

This report, titled “Backing Database System for a Messaging Platform Startup”, was prepared during my 2A work term at Ten Thousand Coffees, where I worked in the Software Engineering department.

Ten Thousand Coffees is startup whose main product is a social networking platform directed towards connecting students with industry professionals. Employed as a member of the core development team for Ten Thousand Coffees, I worked on improving the site and starting experiments over the course of the term. One of the tasks that I was entrusted with was the addition of messaging functionality to Ten Thousand Coffee’s online platform.

From the the initial planning phase of the messaging system until now, there were several changes of requirements, as well as a limited budget for development. One of the fundamental decisions to make during planning phase was the database system to utilizes for this project. This report analyzes two database systems suitable for backing the messaging platform, and identifies the more appropriate alternative for the project.

I would like to thank my co-workers at Ten Thousand Coffees for their continuous guidance during my stay with them. Additionally, I would like to especially thank Elliott Garcea, the lead engineer at Ten Thousand Coffees, for his detailed discussion about our platform’s design decisions with me. Finally, I would like to thank my friend, Raymond Wan, for answering a lot of questions that I have about relational database management systems. I hereby confirm that I have received no help, other than what is mentioned above, in writing this report. I also confirm this report has not been previously submitted for academic credit at this or any other academic institution.

Yours,

---

Clement Hoang, 20531116

## Executive Summary

One of the major decisions in Ten Thousand Coffees' planning phase was the selection of a database management system to use in order to support its messaging service. As a startup, it is imperative that the system of choice allows for rapid prototyping, and has the ability to scale. Some factors that affect the speed of prototyping is the flexibility of schema, and the simplicity of the schema. Additionally, there are multiple ways that databases can scale, which include scaling horizontally, which is the method of adding more machines, or scaling vertically, which is the method of increasing CPU and RAM.

In this report, two popular database management systems, MongoDB, and MySQL, are evaluated against set criteria in order to determine which database is more suited for Ten Thousand Coffee's use case of serving as a datastore for messaging. In terms of schema flexibility, MongoDB was better than MySQL because of its fundamental design philosophies. When scaling, MongoDB is also a better choice because it was designed for horizontal scaling, whereas MySQL was designed when horizontal scaling mattered a lot less. As such, MySQL is more restricted than MongoDB in the context of scaling. Unsurprisingly, MySQL has better query performance than MongoDB based on some conducted experiments where both databases ran on a single machine. Although MySQL wins in this category, it can still be argued that MongoDB would have better performance overall with the right setup and conditions.

Despite MySQL having better query performance than MongoDB in one of the tests conducted, it is not enough to outweigh all the other benefits that MongoDB has to offer. The value of rapid prototyping is worth much higher than a small performance gain, especially at the startup stage. As such, the conclusion is that MongoDB is the best database for use in a messaging service by a startup. Rather, MySQL's use case lies in an established enterprise environment such as a bank, where safe transactions and consistency of data is more important than horizontal scaling and fast prototyping.

# Table of Contents

<b>Executive Summary</b>	<b>iii</b>
<b>Table of Contents</b>	<b>iv</b>
<b>List of Figures</b>	<b>v</b>
<b>List of Tables</b>	<b>vi</b>
<b>1 Introduction</b>	<b>1</b>
<b>2 Background</b>	<b>1</b>
<b>3 Schema Flexibility</b>	<b>2</b>
3.1 MongoDB	3
3.2 MySQL	3
<b>4 Scaling</b>	<b>6</b>
4.1 MongoDB	6
4.2 MySQL	8
<b>5 Performance</b>	<b>9</b>
<b>6 Conclusions</b>	<b>10</b>
<b>7 Recommendations</b>	<b>11</b>
<b>References</b>	<b>13</b>
<b>Acknowledgements</b>	<b>14</b>
<b>The MEAN Stack</b>	<b>A-A.1</b>
<b>Case Study of Schema Migration at Ten Thousand Coffees</b>	<b>A-B.1</b>

## List of Figures

Figure 3.1: A MongoDB user schema for supports group messaging . . . . .	3
Figure 3.2: A MongoDB message schema for supports group messaging . . . . .	4
Figure 3.3: A MongoDB thread schema for supports group messaging . . . . .	4
Figure 3.4: A MySQL schema that supports group messaging [1] . . . . .	5
Figure 4.1: A visual representation of Brewer’s CAP Theorem [1] . . . . .	7
Figure 7.1: A high-level visualization of how the different MEAN components interact .A-A.2	

## List of Tables

Table 5.1: Query Times To Fetch Threads For a Specific User in MongoDB . . . . .	9
Table 5.2: Query Times To Fetch Messages For a Specific Thread in MongoDB . . . . .	10
Table 5.3: Query Times To Fetch Threads For a Specific User in MySQL . . . . .	10
Table 5.4: Query Times To Fetch Messages For a Specific Thread in MySQL . . . . .	10

# 1 Introduction

For several decades, the amount of technological startups has been growing and exploring a vast amount of business models to meet the needs of an ever-changing market. From dating sites to smart-watches, there is a business for almost every idea imaginable. Upon closer inspection of the aforementioned startups, there are noticeably many web applications which are primarily focused on messaging, or, at the very least, include a messaging component. To enable messaging capabilities in a web application, there needs to be a web-server to act as the communication bridge between different users, a database to persist user and message data onto for later usage, and a user interface to allow different people to interact with the web application and message other users. The database, in particular, is very important because it serves as a central store for writing and retrieving user generated content such as message content, chat history logs, and various other datasets for use by the application; without the database, a messaging application simply ceases to function.

However, this is where the startup aspect of Ten Thousand Coffees comes into play. Ten Thousand Coffees is a web application that enables users on the site to connect with one another and invite other parties for coffee chats. As a growing company striving to survive in a fierce environment, the importance of a flexible, scalable, and affordable solution is paramount. This is because startups are limited in budget, and have the need for high growth. In order to be successful, a startup must be efficient in using the resources that it can afford, and make decisions that enable them to scale both quickly, and affordably, so that they could outperform the competition. For the Ten Thousand Coffees platform, many technical decisions were made during the planning phase, and one such decision was to use MongoDB, a new, open-sourced NOSQL database that is based on a document object model that features horizontal scalability as well as a dynamic schema that allows for agile development. MongoDB was chosen over a traditional SQL database such as MySQL. In contrast, SQL databases have been around for much longer than NOSQL databases, and therefore have gained a lot of enterprise users. A SQL database such as MySQL and a NOSQL database such as MongoDB have completely different design principles, and this report will explore why Ten Thousand Coffees chose MongoDB over MySQL.

## 2 Background

As mentioned in the introduction earlier, Ten Thousand Coffees is a social platform whose mission is to connect the youth of today with the leaders of tomorrow over coffee chats. The solution

that the company decided on to approach this problem was a web platform that allows the discovery of possible coffee candidates, with a messaging and scheduling system that allows users to first chat online, before setting up a meeting via the platform and meeting up in real life. The team consists of a business team and a development team, of which contain five and seven employees correspondingly. As a startup with a small team and therefore, limited manpower, there were certain requirements that the chosen database would have to meet in order to be feasible for the platform. These requirements include the ability to prototype rapidly, to be able to scale well, and to be easy-to-use, while still meeting the performance standards for a modern web application. In order to meet these requirements, the team unanimously decided on the MEAN stack, which is a full-stack Javascript framework that became popular over the past few years. For more information on how a web application architected with the MEAN stack works, see Appendix A.

### **3 Schema Flexibility**

The ability to prototype rapidly is critical to a startup because at the early stages of a company, there is no absolute path from beginning to end goal. This means that the requirements will continuously change and adapt based on feedback received from consumers as well as data gathered by analytics. Because of the need to be able to iterate and improve the platform quickly based on changing requirements, the database management system that is used will need to be flexible so that data can quickly and easily be migrated to support the changing requirements. In this section, the abilities of MongoDB and MySQL to meet the requirements of fast prototyping through schema flexibility will be explored.

The criterion of “schema flexibility” refers to the ability of a database to adapt to the changing of the data schema. This is an important measure when comparing the benefits of each database system for a messaging platform startup because through the lifetime of the application, the requirements will continuously evolve to meet the demands of the consumers, and it is also not uncommon for bugs at the schema design level to be introduced where a database schema change is necessary in order to fix it.

For example, an early prototype of a messaging platform may only support conversations in which one user can message only one other user. Obviously, most dedicated messaging platforms will eventually improve its functionality to support features such as multi-user messaging, but as an early iteration, to be able to message only one other user is sufficient. Another scenario in which a schema change may be necessary is when a boolean flag field needs to be added into the schema to contain additional state knowledge will helps to solve certain bugs. In these cases, the database



```

{
  username: {
    type: String,
    index: true
  }
  .
  .
  .
}

```

Figure 3.1: A MongoDB user schema for supports group messaging

management system that is used should be able to migrate or adapt data from the schema of an earlier iteration with a schema of a later iteration with relative ease in order to be a good fit for a startup.

### 3.1 MongoDB

MongoDB is a NOSQL database that differs from a MySQL database in that it uses dynamic schemas. Consequently, the documents stored in a collection do not need to have a uniform data model. This schema flexibility allows for migration methods that are different from the usual script method, in which a large script is executed against the database to modify the data models all the relevant data stored in the database at once. With MongoDB’s dynamic schema, an alternative approach to schema migration, known as “lazy updating” or “schema versioning.” Lazy updating refers to the method of only updating documents as required, having code that supports both the older schema as well as the newer one. This allows the database to slowly migrate without any downtime or a large change all at once, until the older schema version is eventually deprecated and support is dropped. At Ten Thousand Coffees, I led the migration of the older messaging schema into the new architecture in order to add support for group messaging. For the sake of brevity, the details of the migration are omitted from this section. See appendix B for the case study.

### 3.2 MySQL

MySQL is a SQL-compliant database which forces the user to define a structure in terms of tables and columns before data can be stored in it. As such, it demands a long initial setup time. This rigidity comes with benefits such as being more reliable and being more efficient at complex queries [2]. However, while these benefits are great for a more corporate enterprise solution such as banking software, where reliability is more important, they are not as suited for a startup that is

```

{
  content: {
    type: String
  },
  author: {
    type: mongoose.Schema.Types.ObjectId,
    ref: 'User',
    required: true
  },
  date: {
    type: Date,
    default: Date.now
  }
}

```

Figure 3.2: A MongoDB message schema for supports group messaging

```

{
  thread_participants: [{
    user: {
      type: mongoose.Schema.Types.ObjectId,
      ref: 'User',
      required: true,
      index: true
    },
    role: {
      type: String,
      enum: [
        'creator',
        'participant'
      ],
      default: 'participant'
    },
    read: Boolean,
    archived: Boolean,
    updated_at: {
      type: Date,
      default: Date.now
    }
  }],
  messages: [MessageSchema],
  updated_at: {
    type: Date,
    default: Date.now,
    index: true
  },
  created_at: {
    type: Date,
    default: Date.now
  }
}

```

Figure 3.3: A MongoDB thread schema for supports group messaging

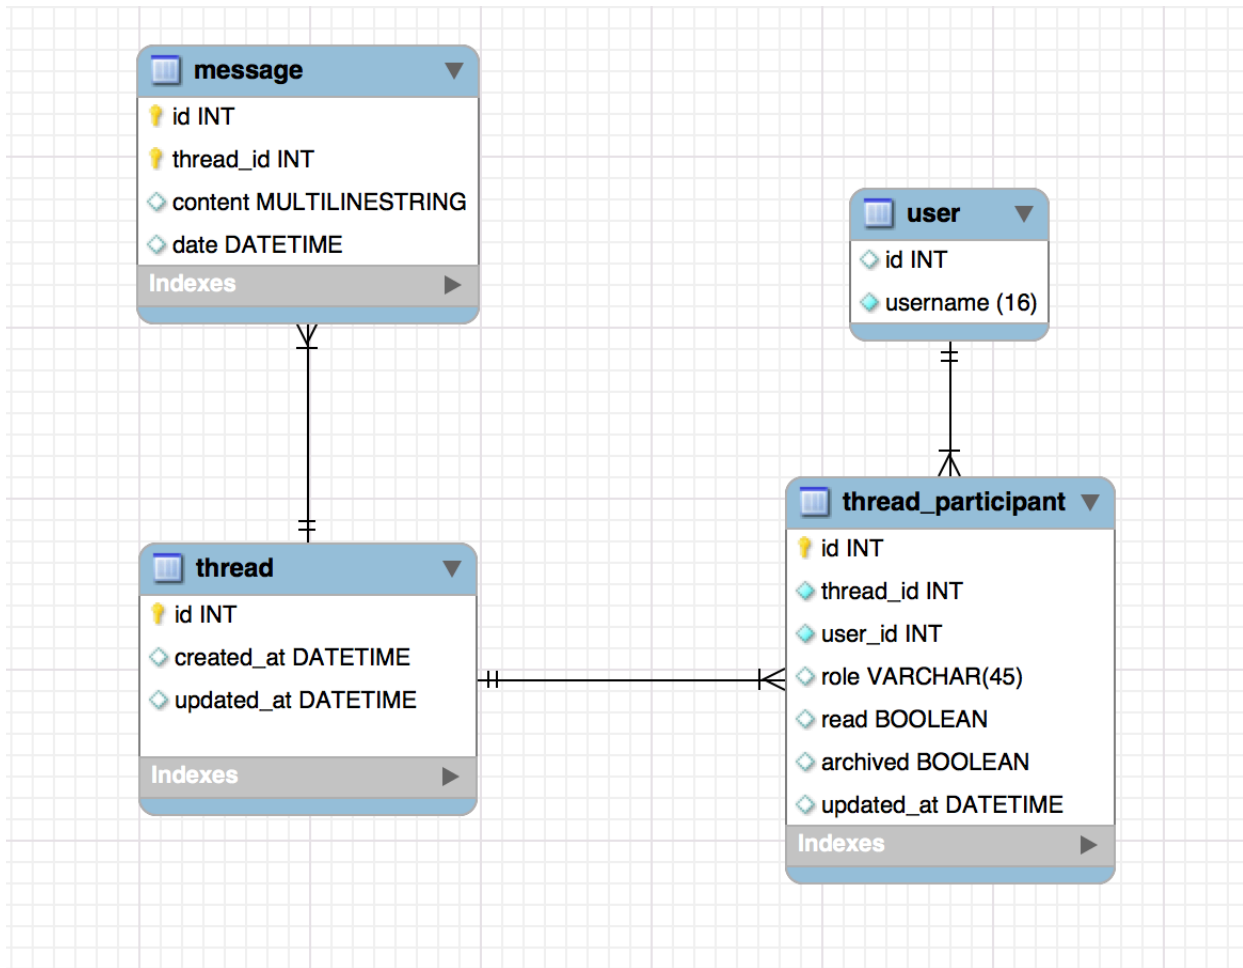


Figure 3.4: A MySQL schema that supports group messaging [1]

focused on a messaging platform. By SQL's nature, it is more complex to set up than MongoDB, and migrations would also require significantly more effort, as the mathematical structure containing tables and rows would require a more revamping than simply adding/removing a field in the case of MongoDB. This is not good for a company that needs fast prototyping. The advantages of rapid prototyping far outweigh the advantages of transaction safety in a startup, and more so in a messaging platform. In the worst case, if MongoDB performance is bad on certain queries, the relevant components can be swapped for something that is more performant such as Redis. Therefore, MongoDB is more suitable than MySQL in terms of schema flexibility.

## 4 Scaling

In a startup, the ability for a database to scale is very important. As a startup often starts off with a limited budget and a small initial consumer base, their product does not need to support high traffic. However, as the company grows and their amount of users rises, then the ability to support the increased traffic is mandatory in order for the company to survive. Different technologies have different techniques to handle specific use cases and so, the characteristics of different databases will vary when they scale. There is a common theorem known as the “CAP” theorem where no system can achieve consistency, high-availability, and partition-tolerance at the same time, and this is made especially apparent when scaling the database management system. A consistent system is one where the system can guarantee that the same state will be reported at all nodes in every subsequent operation until the state is explicitly changed. A system that is highly available is one that is resistant to hardware failure and allows operations to continue even if a node has failed. Finally, a partition-tolerant system is one where the system can be split over several data centers while staying synchronized with one another. [3] For a messaging platform, high availability and partition-tolerance is valued more than consistency.

### 4.1 MongoDB

In terms of the CAP Theorem, MongoDB has very high partition-tolerance. One of the reasons that it is so popular is because of its ability to horizontally scale very well. MongoDB, as a part of its design philosophy, was built for horizontally scaling. That is, unlike most SQL systems that were made years ago when horizontal scaling was not as important, MongoDB was developed in an era where the web is huge and the need to fulfill many concurrent operations at once is very important. MongoDB offers two forms of horizontal scaling: replication and sharding.

Replication provides redundancy by copying data on different database servers. This provides benefits such as protecting a database from the loss of a single server, and with additional copies of the data, one database can also be dedicated to disaster recovery, reporting, or backup. Additionally, replication can be used to increase read capacity, as it allows clients to write and read to different servers. Copies can also be strategically placed in different data centers to increase locality and availability of data for distributed applications [4].

MongoDB also allows horizontally scaling via sharding. Sharding is a method for storing data across multiple machines. MongoDB uses sharding to support deployments with very large data sets and high throughput operations. Database systems with large data sets and high throughput

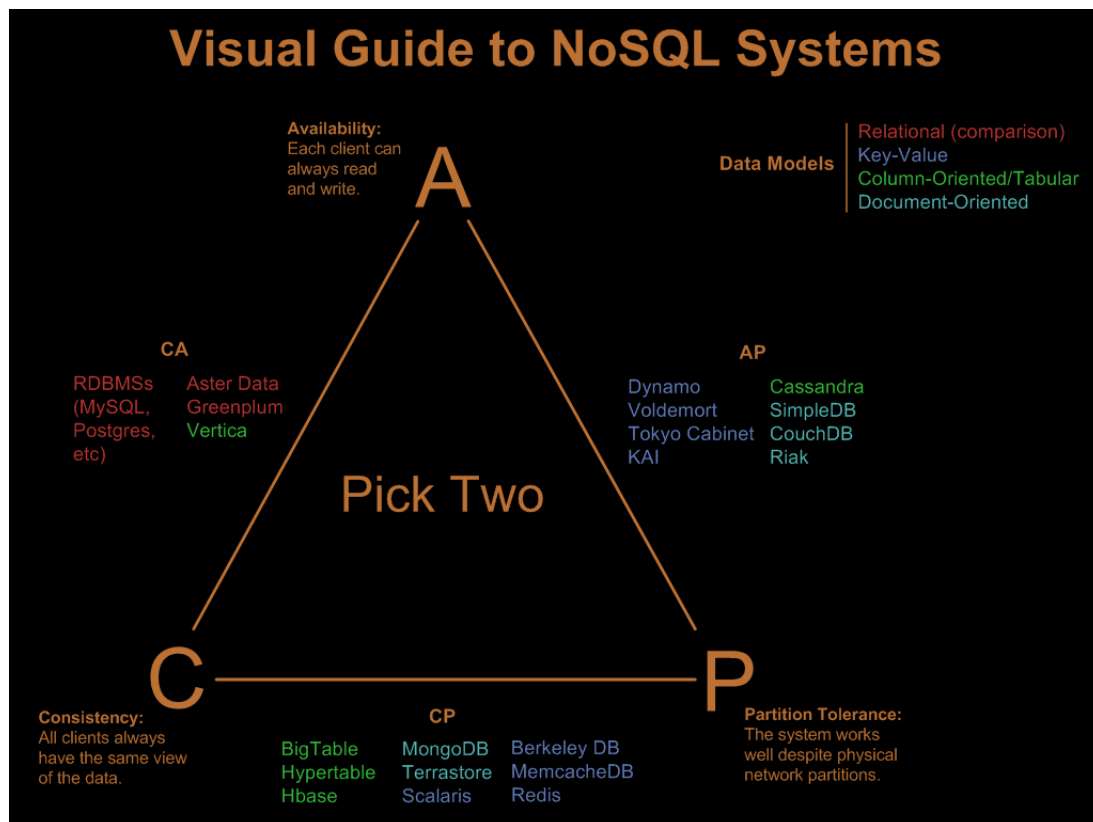


Figure 4.1: A visual representation of Brewer's CAP Theorem [1]

applications can challenge the capacity of a single server, since high query rates can exhaust the CPU capacity of the server, and larger data sets easily exceed the storage capacity of a single machine. Finally, working set sizes larger than the system's RAM stress the I/O capacity of disk drives. Sharding, by contrast, divides the data set and distributes the data over multiple servers, or shards. Each shard is an independent database, and collectively, the shards make up a single logical database [5].

As for vertical scaling, more CPU and high-capacity storage resources can simply be added to a database system in order to make it perform better. However, it has a practical limit and the becomes exponentially expensive, so horizontal scaling is much more efficient than vertical scaling for MongoDB.

With MongoDB's scaling capabilities, tens of thousands of organizations have used MongoDB in their high performance systems. In fact, over 30 of the Fortune 100 have grown from single server deployments to clusters with over 1,000 nodes, delivering millions of operations per second on over 100 billion documents and petabytes of data [6].

## 4.2 MySQL

In terms of the CAP Theorem, MySQL falls in the consistent and high availability categories. As mentioned earlier, this means that MySQL's strengths lie in its ability to report the same state to all users accessing the database, and its resilience to downtime after operations are completed. Additionally, it has the option, like all other databases, to vertically scale by increasing memory and CPU. However, because of MySQL's relational database structure that contains tables, it is harder for MySQL to scale horizontally [7]. That is, even with sharded data on MySQL, queries with joins will still be quite inefficient because of the latency from when the query is executed, then the server parsing, optimizing, finding an execution plan, and executing it, going to all the different tables across different data stores to retrieve the queried data. Despite this, there are still some horizontal scaling solutions for MySQL which have been improving over the year from the use of clever algorithms, such as ScaleBase and MySQL Cluster.

The analysis of the scaling capabilities of both MongoDB and MySQL shows that although MySQL has some capability to scale horizontally, it requires much more effort and is not as good as MongoDB's out of the box sharding support. Additionally, both databases can accommodate vertical scaling, so the result is that MongoDB wins in the category of scalability.

## 5 Performance

When comparing the database's performances against one another, there are a myriad of system setups that can be used in order to conduct the test. However, this falls under the category of scalability that was mentioned earlier in this report. Instead, this section will evaluate the performance of MongoDB and MySQL based on my local machine with the same default setup. These tests do not aim to measure the pure performance metrics of the two databases, but attempt to eliminate as much bias as possible by performance common queries that would be used in a group messaging application, and recording the metrics for the relevant querying needs.

These querying needs include fetching threads for a specific user, and fetching messages in a specific thread. The tests for both MongoDB and MySQL are conducted using a 2.5 GHz Unix machine with 16 GB of RAM. The test data that is used contains 20,000 users, with an average of 20 threads of roughly 50 messages each.

Table 5.1: Query Times To Fetch Threads For a Specific User in MongoDB

Trial #	Query Time
1	311 ms
2	298 ms
3	305 ms
4	299 ms
5	318 ms
6	294 ms
7	301 ms
8	307 ms
9	315 ms
10	309 ms

From the results of the ten-trial experiments that I ran, it can be inferred that MySQL's query times for most of the fetch operations are slightly quicker than those of MongoDB. However, they are still very comparable, and with other techniques, the differences can be deemed negligible. For example, MongoDB can improve its lookup speed by indexing key fields. In addition, as the data models become more complex, MongoDB has the superior lookup model since it stores data as document models rather than having to look up rows in different tables. Finally, MongoDB has the advantage of horizontal scaling, as covered in a previous section, so its performance on a single machine is not critical. The conclusion in here is that MySQL has slightly faster message and thread fetch times than MongoDB on a single machine.

Table 5.2: Query Times To Fetch Messages For a Specific Thread in MongoDB

Trial #	Query Time
1	250 ms
2	257 ms
3	242 ms
4	239 ms
5	245 ms
6	242 ms
7	248 ms
8	240 ms
9	238 ms
10	251 ms

Table 5.3: Query Times To Fetch Threads For a Specific User in MySQL

Trial #	Query Time
1	122 ms
2	137 ms
3	134 ms
4	130 ms
5	128 ms
6	127 ms
7	133 ms
8	131 ms
9	129 ms
10	135 ms

Table 5.4: Query Times To Fetch Messages For a Specific Thread in MySQL

Trial #	Query Time
1	110 ms
2	112 ms
3	109 ms
4	116 ms
5	108 ms
6	107 ms
7	113 ms
8	111 ms
9	110 ms
10	106 ms

## 6 Conclusions

In this report, MongoDB and MySQL have been evaluated on their schema flexibility, scalability, and performance, in order to assess their suitability for a startup focused on a messaging platform.



In schema flexibility and scalability, MongoDB trumped against MySQL, but it lost to MySQL when doing query benchmarking.

Schema flexibility is an important factor for a social startup because the schemas evolve very quickly in early stages; in the prototyping stage, the schema undergoes constant changes to meet the ever-changing demands. With a qualitative analysis of MongoDB's design philosophies as well as a qualitative analysis of MySQL's design philosophies, it is evident that MongoDB's design decision of a dynamic schema allows it to be far more flexible than the mathematical table structure of MySQL. This allows for faster and easier schema iterations and versions.

In terms of scalability, both MySQL and MongoDB were able to vertically scale. However, MongoDB is able to scale horizontally better than MySQL is able to. Although MySQL has tools that allow it to scale horizontally, that's not what it was originally designed for, and as such, it lags behind MongoDB in horizontal scalability. In the modern technological era, horizontal scalability is much more important than vertical scalability because it means less development effort put into the initial iteration, but allows an easy way to scale. In addition, it is less expensive than vertical scaling.

MySQL has a slight edge against MongoDB in common queries for group messaging on a single machine. However, this slight performance edge is not enough to offset the benefits of rapid prototyping and ease of development that MongoDB offers. The performance of MongoDB and MySQL are comparable, and there are many ways to improve MongoDB's performance, which was bottlenecked, since the experiments were performed on a single machine. Regardless, MySQL wins in this category.

Having a better result than MySQL in two out of the three criteria, MongoDB is the database management system of choice when building a messaging platform for a startup. It allows for fast prototyping, and is great at scalability, and the only expensive is a very slight performance reduction, which can be minimized with many techniques in the later stages of the startup.

## **7 Recommendations**

In the future, when MongoDB is selected as the database management system for a messaging platform startup, the allocation of human resources and other technologies should be set according to the benefits that MongoDB brings. For example, MongoDB uses JSON-like documents as its data model. This means that it would fit best with Javascript technologies such as Node.js. If this

recommendation is followed, then it would mean that a majority of the technology stack would be in Javascript. This means that software developers who are experienced in Javascript should be the better hire, and Javascript should be learned by both front-end developers and back-end developers.

Also, since MongoDB has a very simple document structure, the learning curve for it is quite low. The structure does not quickly become unmanageable like SQL where it is common to see projects with hundreds of tables. With a SQL database, it is a necessity to hire a database administrator in order to make design decisions and monitor the database. However, because of MongoDB's developer-friendliness, every software developer on the development team should be able to utilize the database with relative ease, and a specialized database administrator is not necessary.

It is also notable that this report is based strictly on the use case of a messaging platform for a startup. For example, Facebook has a messaging service, but its architecture will deviate greatly from the infrastructure that is outlined in this report. This is because Facebook has the technology and manpower to have a customized solution which handles requests on a scale that is much larger than what startups have to deal with. It would use an in-memory database like Redis for instant messages, which would synchronize to a more robust backing database; it is a custom solution that startups would not have to deal with for a long time. Such a large scale is not important for startups. Rather, what is important is fast prototyping and scalability.

## References

- [1] N. Hurst. (2010, Mar.) Visual guide to nosql systems. [Online]. Available: <http://blog.nahurst.com/visual-guide-to-nosql-systems>
- [2] O. Tezer. (2014, Feb.) Understanding sql and nosql databases and different database models. [Online]. Available: <https://www.digitalocean.com/community/tutorials/understanding-sql-and-nosql-databases-and-different-database-models>
- [3] R. Tharakan. (2010, Feb.) Brewers cap theorem on distributed systems. [Online]. Available: <http://www.royans.net/wp/2010/02/14/brewers-cap-theorem-on-distributed-systems/>
- [4] (2015) Replication introduction. [Online]. Available: <http://docs.mongodb.org/manual/core/replication-introduction/>
- [5] (2015) Sharding introduction. [Online]. Available: <http://docs.mongodb.org/manual/core/sharding-introduction/>
- [6] (2015) Mongodb at scale. [Online]. Available: <http://www.mongodb.com/mongodb-scale>
- [7] P. Campaniello. (2011, Jun.) Mysql query engine scalability issues. [Online]. Available: <https://www.scalebase.com/mysql-query-engine-scalability-issues/>

## **Acknowledgements**

I would like to thank my co-workers at Ten Thousand Coffees for their continuous guidance during my stay with them. Additionally, I would like to especially thank Elliott Garcea, the lead engineer at Ten Thousand Coffees, for his detailed discussion about our platform's design decisions with me. Finally, I would like to thank my friend, Raymond Wan, for answering a lot of questions that I have about relational database management systems.

## Appendix A The MEAN Stack

The MEAN stack is a full-stack Javascript framework for web applications. It differs from other web frameworks such as Ruby on Rails, or Python with Django in that it is a monoglot framework and only uses Javascript for both its front-end and its backend logic. Because the MEAN stack is fully Javascript, that means most web developers will be able to learn how to program with the MEAN stack with relative ease, as Javascript is virtually the only front-end programming language that exists for web development, and they will have had at least some experience with it.

MEAN is an acronym that stands for MongoDB, Express.js, Angular.js, and Node.js. This section of the appendix will explain each component of the MEAN stack, and how they interact with each other.

MongoDB is a document-oriented database management system. Rather than the relational tables, it contains JSON-like documents. One major difference from RDBMS is that it doesn't support joins nor transactions, but instead it nests documents.

Express is a minimal and flexible Node.js web application framework that provides a robust set of features for web applications. It provides a nice interface for Node to route client requests, and it also features a middleware system that requests can bubble through.

Angular.js is a structural MV\* framework for dynamic web applications that have the need for two-way data binding. It extends HTML's syntax to express application components clearly and concisely. It also uses dependency injection to make code cleaner and more modular.

Node.js is a platform built on Chrome's Javascript runtime that uses an event-driven, non-blocking I/O model. Its use cases are primarily for data-intensive real-time applications that require concurrency. On the contrary, it is not a good choice for projects that require CPU-intensive operations.

With the MEAN architecture, Node with Express is the event-driven server that handles user requests as well as API calls from the client. The user would be interacting with the Angular site that is served statically via the Express server, and their interaction with the site would lead to actions being triggered. These actions will interact with the application's API through HTTP requests. Through the web application's API, updates, retrievals, and posts would be made to the fields in the MongoDB database.




Figure 7.1: A high-level visualization of how the different MEAN components interact

## **Appendix B Case Study of Schema Migration at Ten Thousand Coffees**

During my co-op term at Ten Thousand Coffees, another co-op student implemented a messaging service into the platform that supported one-to-one conversations. He used fields such as “receiver”, “sender”, “receiver\_updated”, and “sender\_updated” inside the thread schema. However, a schema like is evidently unable to support more than two users per thread. The only way to support a variable amount of users in a thread was to include an array of thread participants in each thread. In my past experience using a relational database, migrating the old data to keep the history of old conversations was going to require significant effort. Although, I later discovered that the task was not too difficult. My co-worker advised that the best practice was to use lazy updating.

For the entire migration, I only had to make a few changes. In the thread schema, I added a “thread\_participants” array that would contain the all the users in the chat, while leaving the older “receiver” and “sender” fields intact. I then added a virtual method that simply reports whether a document has an older version of the schema, or the newer version of the schema. Using this method, I changed the code on the client to handle both versions, but once an update is performed, to convert the older version to the newer one. This way, migrating was relatively pain-free.